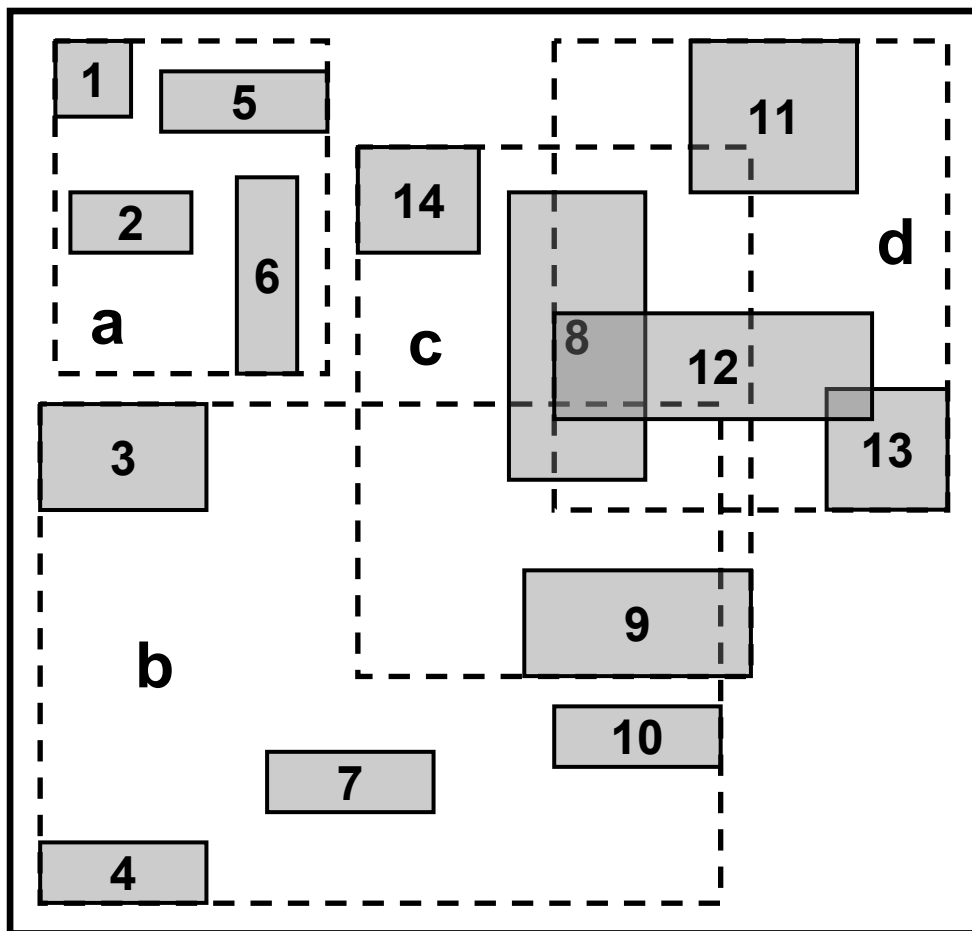



R-Tree

- An R-tree is a depth-balanced tree
 - Each node corresponds to a disk page
 - Leaf node: an array of leaf entries
 - A leaf entry: (mbb, oid)
 - Non-leaf node: an array of node entries
 - A node entry: (dr, nodeid)



$m=2, M=4$



Properties

- The number of entries of a node (except for the root) in the tree is between m and M where $m \in [0, M/2]$
 - M : the maximum number of entries in a node, may differ for leaf and non-leaf nodes
 $M = \lfloor \text{size}(P) / \text{size}(E) \rfloor$ P : disk page E : entry
 - The root has at least 2 entries unless it is a leaf
- All leaf nodes are at the same level
- An R-tree of depth d indexes at least m^{d+1} objects and at most M^{d+1} objects, in other words, $\lfloor \log_M N - 1 \rfloor \leq d \leq \lfloor \log_m N - 1 \rfloor$

Search with R-tree

- Given a point q , find all mbbs containing q
- A recursive process starting from the root

$result = \emptyset$

For a node N

if N is a leaf node, then $result = result \cup \{N\}$

else // N is a non-leaf node

for each child N' of N

if the rectangle of N' contains q

then recursively search N'

Time complexity of search

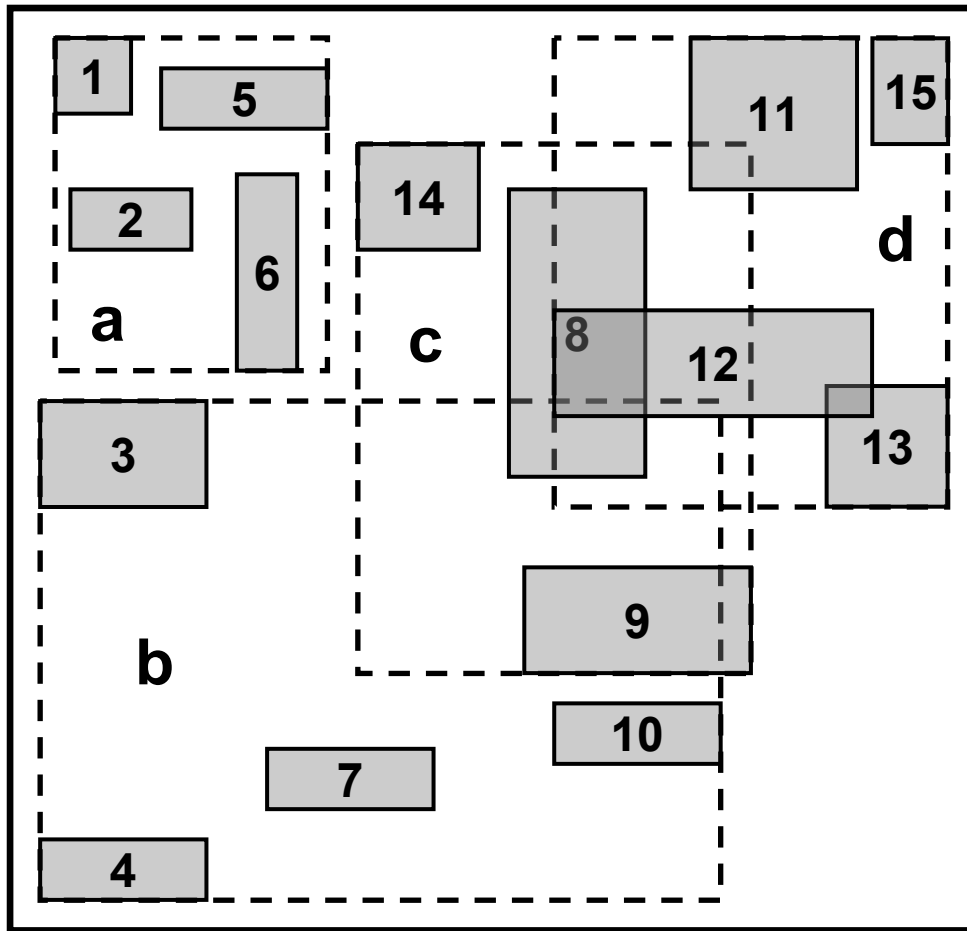
- If mbbs do not overlap on q , the complexity is $O(\log_m N)$.
- If mbbs overlap on q , it may not be logarithmic, in the worst case when all mbbs overlap on q , it is $O(N)$.

Insertion – choose a leaf node

- Traverse the R-tree top-down, starting from the root, at each level
 - If there is a node whose directory rectangle contains the mbb to be inserted, then search the subtree
 - Else choose a node such that the enlargement of its directory rectangle is minimal, then search the subtree
 - If more than one node satisfy this, choose the one with smallest area,
- Repeat until a leaf node is reached


Insertion – insert into the leaf node


- If the leaf node is not full, an entry [mbb, oid] is inserted
- Else // the leaf node is full
 - Split the leaf node
 - Update the directory rectangles of the ancestor nodes if necessary



Insert object 15

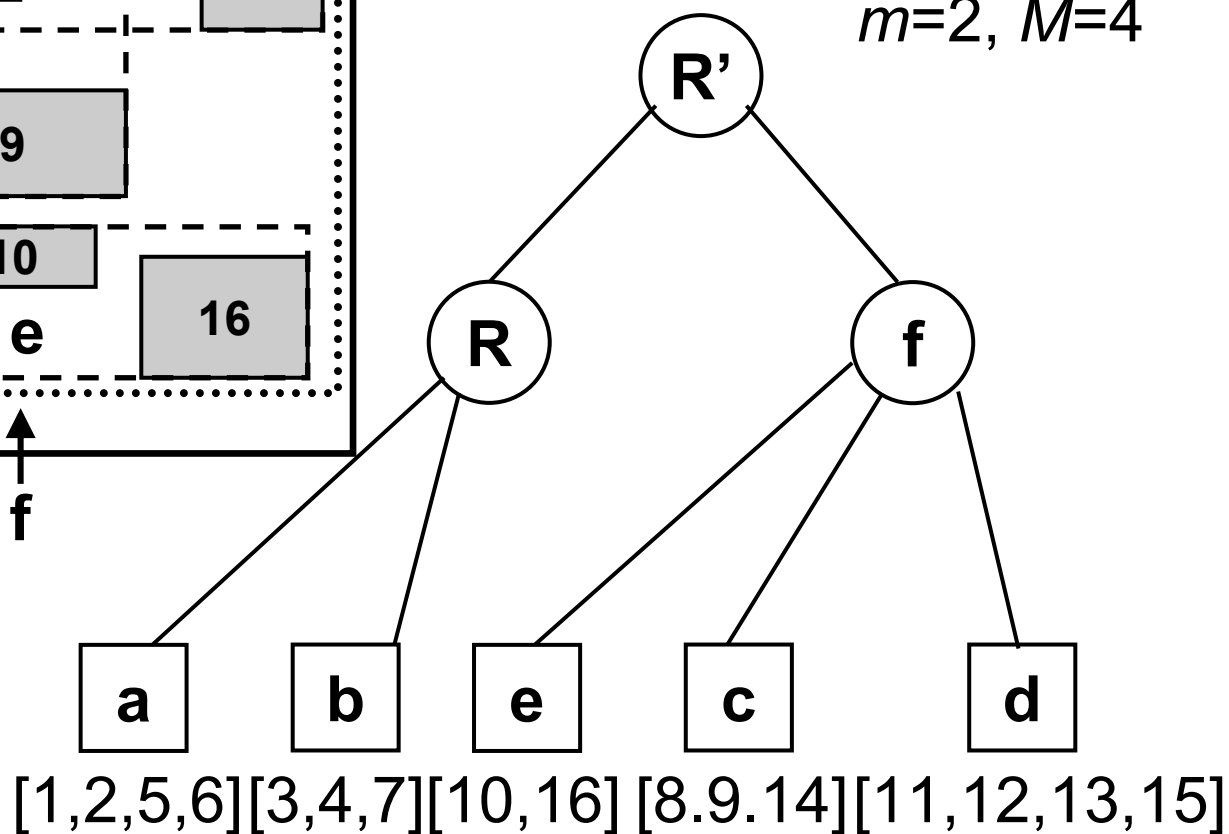
$m=2, M=4$







Insert object 16

$m=2, M=4$



Split - goal

- The leaf node has M entries, and one new entry to be inserted, how to partition the $M+1$ mbbs into two nodes, such that
 - 1. The total area of the two nodes is minimized
 - 2. The overlapping of the two nodes is minimized
- Sometimes the two goals are conflicting
 - Using 1 as the primary goal



Split - solution

- Optimal solution: check every possible partition, complexity $O(2^{M+1})$
- A quadratic algorithm:
 - Pick two “seed” entries e_1 and e_2 far from each other, that is to maximize $\text{area}(\text{mbb}(e_1, e_2)) - \text{area}(e_1) - \text{area}(e_2)$
here $\text{mbb}(e_1, e_2)$ is the mbb containing both e_1 and e_2 , complexity $O((M+1)^2)$
 - Insert the remaining $(M-1)$ entries into the two groups

Quadratic split cont.

- A greedy method
- At each time, find an entry e such that e expands a group with the minimum area, if tie
 - Choose the group of small area
 - Choose the group of fewer elements
- Repeat until no entry left or one group has $(M-m+1)$ entries, all remaining entries go to another group
- If the parent is also full, split the parent too. The recursive adjustment happens bottom-up until the tree satisfies the properties required. This can be up to the root.